

ABDULLAH SHAFIQUE

AI Engineer | Generative AI • LLM Apps • RAG • AI Agents

CONTACT

Location
Sahiwal, Punjab, Pakistan

Phone
+92 309 950 8436

Email
abdullahshafique2019@gmail.com

LinkedIn
linkedin.com/in/aadi-abdullah

Portfolio
aadi-abdullah.github.io

GitHub
github.com/aadi-abdullah

TECH STACK

- Python
- LangChain
- LangGraph
- FastAPI
- Groq
- FAISS
- ChromaDB
- Streamlit
- Scikit-learn
- XGBoost
- Docker
- GitHub Actions
- Pandas
- Prompt Engineering

CERTIFICATIONS

- AI Agent Developer Specialization
Vanderbilt University
- Microsoft Python Development
Specialization
- Adobe Graphic Designer
Specialization

EDUCATION

BSc Software Engineering
Riphah International University
GPA: 3.99 / 4.0 • 2024 – 2028

LANGUAGES

English (Professional)
Urdu (Native)

SUMMARY

AI Engineer with 6 years of graphic design experience. I build autonomous agents, RAG pipelines, and LLM-powered applications using Python, LangChain, LangGraph, and FastAPI. Design background means I optimise for usefulness — every system I build is made to be understood, used, and trusted by real people.

PROJECTS

AI Research Agent

Live Demo → [aadi-research-agent.streamlit.app](#)

LangChain · Groq Llama 3.3 70B · Tavily · Streamlit

70% faster research

Llama 3.3 70B

Real-time web search

- Autonomous agent that searches 5 live sources via Tavily, synthesises findings with Groq, and exports a structured PDF report in under 60 seconds.
- LangChain orchestration with Search Agent → Summarizer chain → Report Agent → PDF Exporter pipeline.

PDF Knowledge Assistant

Live → [pdf-knowledge-assistant.vercel.app](#)

FastAPI · React 18 · FAISS · Groq Llama 3.1 · Sentence-Transformers

Recall@5: 0.91

<3s answers

Page-level citations

- RAG system — upload any PDF, ask questions in plain English, receive grounded answers with page-level source citations.
- FAISS vector store runs locally (no cloud DB). React frontend on Vercel, FastAPI backend on Hugging Face Spaces.

House Price Prediction API

[github.com/aadi-abdullah/house-price-prediction-api](#)

FastAPI · XGBoost · Scikit-learn · Docker · Uvicorn

RMSE: 0.48

p95: 210ms

120 req/s

- End-to-end ML pipeline: data preprocessing → model training (XGBoost, RandomForest, Linear) → best-RMSE selection → FastAPI REST endpoint.
- Pydantic request validation, Docker + Unicorn deployment, confidence interval in every response.

LLM Evaluation Framework

[github.com/aadi-abdullah/llm-eval-framework](#)

Python · ROUGE · BERTScore · GitHub Actions CI

50 test prompts

2 regressions caught

- Automated testing framework for LLM outputs using ROUGE-L, BERTScore, and a custom NLI faithfulness checker.
- Integrated with GitHub Actions — posts score regression alerts as PR comments before deployment.

EXPERIENCE

AI Engineer — Independent Projects

2025 – Present

Self-Directed

- Building production AI systems: autonomous agents, RAG pipelines, ML APIs, and LLM evaluation frameworks.
- Focused on measurable outcomes — every project ships with evaluation metrics, latency data, and failure handling.

Freelance Graphic Designer

Jan 2018 – Oct 2024 (6 yrs 10 mo)

Self-Employed

- Designed logos, brand identities, and digital marketing visuals for startups and creators across Pakistan.
- Product thinking and user-focused design now applied directly to AI interface and system design.